

---

---

# Introduction to Kubernetes!

— Savitha Raghunathan —

---

---

# Introduction

Savitha Raghunathan

- Senior Software Engineer @ Red Hat
- Kubernetes Contributor
- CNCF Ambassador

# Agenda

Part 1: Containers & Docker

Part 2: Kubernetes

Part 3: Kubernetes Architecture

Part 4: Kubernetes Components

Part 5: Hands-On Demo

Part 6: Wrap-Up and Q&A

# Part 1: Containers & Docker

# Traditional Virtual Machines and Local Setups

Traditional virtual machines (VMs) virtualize an entire operating system, and running applications locally on a single OS can lead to conflicts.

Disadvantages:

- Heavy Resource Usage: VMs replicate the whole operating system, consuming substantial resources (RAM, CPU) on each VM instance.
- Lack of Consistency: Running apps locally can result in version conflicts (dependencies, environment variables)
- Inefficient Scaling: Spinning up a new VM is slower and requires more storage

# What are Containers?

Containers package an application with its dependencies, letting it run consistently across any environment.

Key Points:

- Containers occupy a smaller footprint compared traditional VMs.
- Containers allow consistency by packaging everything the application needs.
- Containers can be moved between environments without issues.

# Why Use Containers?

Benefits:

- Portability: Containers can work across all computers and servers.
- Efficiency: Containers only carry what's needed for the app, making them smaller and faster than full virtual machines.
- Resource Isolation: Containers are independent, preventing conflicts.

# Introduction to Docker

Docker is the tool used to create and manage containers.


Key Commands:

**docker build**: Assembles the container image.

**docker run**: Runs the container.

**docker pull**: Downloads Docker images from the Docker Hub.

**docker ps**: Lists running containers, providing information about their status and configurations.





# Running a Container with Docker

Run an Nginx container:

```
docker run --name my-nginx -p 8080:80 nginx
```

Open a browser and visit localhost:8080 to see the Nginx welcome page.

# Running a Container with Docker

```
docker run --name my-nginx -p 8080:80 nginx
```

- `--name my-nginx`: Names the container instance `my-nginx`.
- `-p 8080:80`: Maps port 80 in the container to port 8080 on the host, making the application accessible locally.
- `nginx`: Specifies the image to use (in this case, the latest Nginx image from Docker Hub).

# Part 2: Kubernetes

# Why Do We Need Kubernetes?

Complexity at Scale: Managing multiple containers across distributed environments is resource-intensive and complex.

- Scaling Needs: Adjusting container numbers manually to match traffic spikes or dips requires significant oversight.
- Configuration Drift: Without consistent orchestration, container configurations can drift across environments, causing reliability issues.

# Why Do We Need Kubernetes?

High Availability & Resilience:

- Fault Tolerance: Detecting and recovering from container or node failures promptly is crucial to avoid downtime.
- Service Discovery: Ensuring containers can locate each other (especially when they're dynamically scheduled across nodes) is challenging without orchestration.

# Why Do We Need Kubernetes?

Load Distribution:

- Traffic Management: Balancing incoming traffic across containers to avoid overload or resource wastage is difficult to manage manually.

# Kubernetes Overview

An orchestration platform that automates the deployment, scaling, and management of containerized applications.

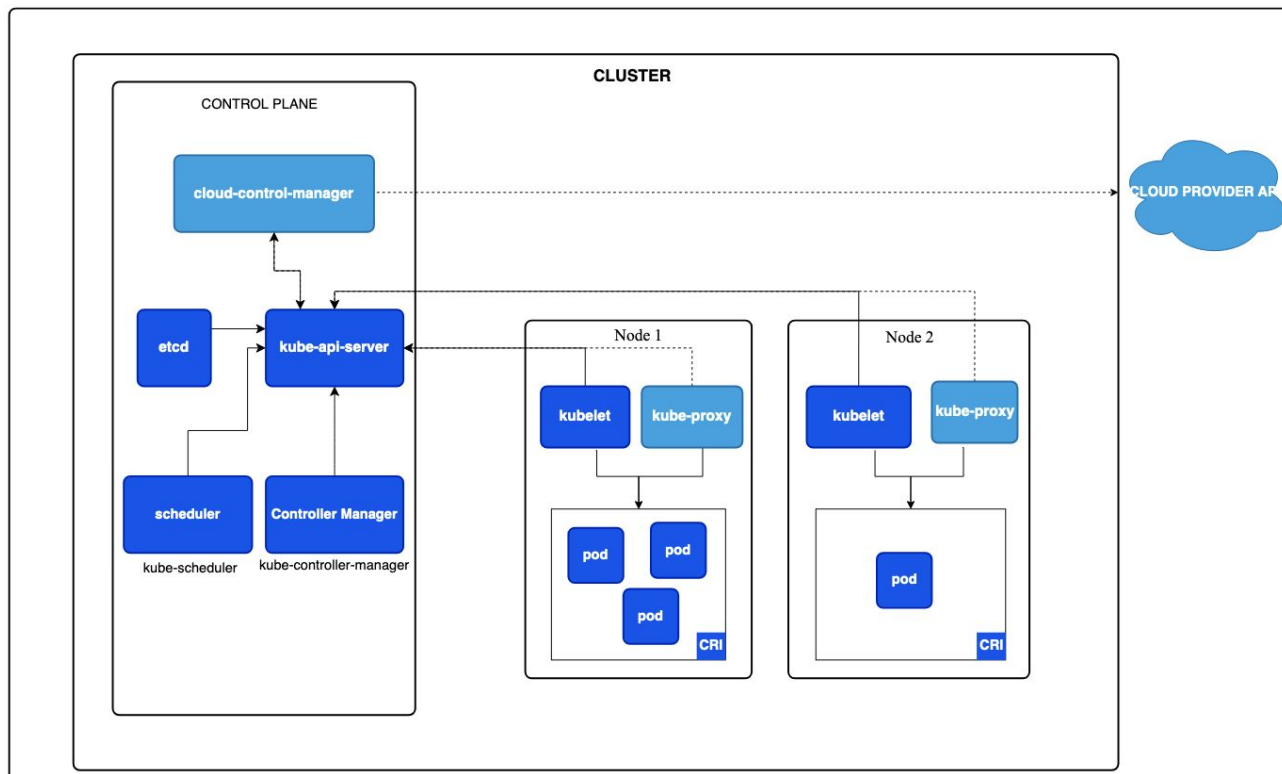
Main Responsibilities:

- **Scaling:** Automatically adjusts the number of running containers based on real-time demand, adding or removing instances as needed.
- **Load Balancing:** Distributes incoming network traffic evenly across containers, ensuring efficient resource utilization and application performance.
- **Self-Healing:** Monitors container health, automatically replacing or restarting containers that fail, maintaining application availability and reliability.

# Part 3: Kubernetes Architecture



# Kubernetes Architecture



Pic credits: <https://kubernetes.io/docs/concepts/architecture/>

# Kubernetes Control Plane

Control Plane: The core management layer of Kubernetes that maintains the desired state of the cluster.

- API Server: The central component that exposes the Kubernetes API and serves as the communication point for all Kubernetes operations.
- Scheduler: Assigns Pods to available nodes based on resource requirements and policies.
- Controller Manager: Ensures the actual state matches the desired state, handling tasks such as maintaining replicas and managing node lifecycles.
- etcd: A highly available key-value store that stores all cluster data, configurations, and states.

# API Server

The API Server (kube-apiserver) is the main communication point for the entire Kubernetes cluster.

- It handles all REST API requests from users, CLI tools (kubectl), and other Kubernetes components, serving as the front-end interface for interacting with the cluster.
- Responsibilities:
  - Validates and processes REST requests.
  - Acts as the gatekeeper for all operations (CRUD) on Kubernetes objects.
- Security: It enforces authentication, authorization, and admission control policies, ensuring secure interactions with the cluster.

# Scheduler

The Scheduler (kube-scheduler) is responsible for assigning new Pods to nodes within the cluster based on resource availability and scheduling policies.

Responsibilities:

- Monitors unscheduled Pods and assigns them to nodes with sufficient resources (CPU, memory).
- Applies scheduling rules, such as affinity and anti-affinity (e.g., ensuring two specific Pods are scheduled together or kept separate).
- Optimizes workload distribution to maintain efficient resource usage and avoid overloading nodes.

Decision Criteria: The Scheduler considers factors like resource requirements, node capacity, and existing workloads to make scheduling decisions.

# Controller Manager

The Controller Manager (kube-controller-manager) oversees various controller processes that regulate and reconcile the state of Kubernetes objects.

Responsibilities:

- Node Controller: Monitors node health and reacts if nodes become unresponsive.
- Replication Controller: Ensures the specified number of replicas are running for each application.
- Endpoint Controller: Manages endpoint objects and links Services to Pods.
- Service Account & Token Controller: Manages default accounts and API tokens for newly created namespaces.

Reconciliation Loop: The Controller Manager constantly compares the desired state with the current state and makes adjustments as needed, following a control loop process.

# etcd

etcd is a distributed key-value store that serves as Kubernetes' primary data store, holding all cluster configuration data, metadata, and the state of the entire system.

Responsibilities:

- Stores configuration data such as API objects, deployment details, and cluster states.
- Provides high availability with data replication, ensuring that the cluster's state is accessible across multiple master nodes.

Data Consistency: Since Kubernetes heavily relies on etcd for storing the cluster's desired state, it's critical to back up etcd regularly to prevent data loss.

# Kubernetes Worker Nodes

Worker Nodes: Nodes that run containerized applications and provide compute capacity to the cluster.

- Kubelet: An agent that manages Pod operations on each node, ensuring containers are running and healthy.
- Kube-proxy: A network proxy that manages network communication, routing traffic to the appropriate Pods within the cluster.

# Kubelet

The Kubelet (kubelet) is an agent that runs on every worker node, acting as the primary node-level controller.

Responsibilities:

- Receives Pod definitions from the API Server and ensures they are running as specified.
- Monitors container health and reports node status back to the control plane.
- Works with container runtimes (e.g., Docker, containerd) to manage container lifecycle events like starting, stopping, and restarting containers.

Node Monitoring: The Kubelet continually monitors the state of Pods and containers, sending status updates back to the API Server.



# Kube-proxy

The Kube-proxy is a network proxy that runs on each worker node, managing networking rules to allow seamless communication between services and Pods within the cluster.

Responsibilities:

- Configures IP tables or IPVS rules to route traffic to the correct Pods based on Service endpoints.
- Balances traffic across multiple Pods within a Service, enabling load balancing and ensuring even resource utilization.
- Facilitates communication between internal and external clients by handling NodePort or LoadBalancer configurations, making services accessible outside the cluster.

Traffic Routing: The Kube-proxy is essential for forwarding and routing traffic to the appropriate endpoints, ensuring Kubernetes applications remain accessible.

# Container Runtime

The container runtime is the underlying software responsible for running containerized applications on each worker node. Common runtimes include Docker, containerd, and CRI-O.

Responsibilities:

- Pulls container images from repositories, creates containers, and executes them within Pods.
- Provides the environment required for running applications within isolated containers.

CRI (Container Runtime Interface): Kubernetes supports multiple container runtimes through the CRI, enabling flexibility and compatibility with different runtime technologies.

# Part 4. Kubernetes Concepts

# Pods - The Basic Unit

A Pod is the smallest deployable unit in Kubernetes, consisting of one or more tightly coupled containers that share storage, network, and a single IP address.

Key Features:

- Shared Resources: Containers within a Pod share network interfaces and storage volumes.
- Use Case: Often used to run a single container, but can contain multiple containers that need to work closely together.

*Example:* A web application Pod might contain both the main web server container and a sidecar container for logging or monitoring.

# Deployments

A Deployment is a Kubernetes object that defines and manages the desired state for a set of Pods, specifying the number of replicas, the container image version, and the update strategy.

Key Features:

- **Scaling:** Automatically increases or decreases the number of replicas based on demand.
- **Rolling Updates:** Ensures zero-downtime updates by gradually replacing old versions of Pods with new ones.
- **Self-Healing:** Automatically replaces failed Pods to maintain the specified replica count.

*Example:* A Deployment for a front-end application might specify three replicas, ensuring that three instances of the application are always available.

# Services

A Service is a Kubernetes resource that defines a stable IP and DNS name for a set of Pods, enabling communication between different parts of an application or external access to Pods.

Types of Services:

- ClusterIP: Exposes the Service within the cluster, allowing communication between internal Pods.
- NodePort: Exposes the Service on a static port on each node, allowing external access.
- LoadBalancer: Provisions an external load balancer, useful for cloud-based applications.

*Example:* A Service can expose a group of web server Pods to other Pods within the cluster or to external users, ensuring continuous access even if individual Pods are replaced or rescheduled.

# Part 5: Kubernetes in Action

# Demo

- Deploy a cluster using Kind
- Create an nginx webserver deployment
- Expose the deployment using a Service object
- Scale the Deployment
- Delete a pod and watch the deployment



# Part 6: Resources and Q/A

# Recap & Further Resources

- Containers: Provide consistent, portable application packaging by bundling code and dependencies. They enable reliable application deployment across diverse environments.
- Kubernetes: An orchestration platform that automates container management, ensuring scalability, fault tolerance, and efficient resource allocation for containerized applications.
  - Kind: A lightweight tool for running Kubernetes clusters locally, ideal for development, testing, and hands-on learning with Kubernetes.

# Recap & Further Resources

- Docker Tutorials: Resources for understanding Docker and containerization fundamentals – <https://www.docker.com/101-tutorial/>
- Kubernetes Documentation: In-depth guides on Kubernetes concepts, architecture, and components – [kubernetes.io/docs](https://kubernetes.io/docs)
- Kind Documentation: Learn more about setting up and managing local clusters – <https://kind.sigs.k8s.io/docs/user/quick-start/>

# Recap & Further Resources

- Link to slides: <https://github.com/savitharaghunathan/101-kubernetes>



**THANK YOU :) :)**