

---

---

# Introduction to Kubernetes!

— Savitha Raghunathan —

---

---

# Introduction

Savitha Raghunathan

- Senior Software Engineer @ Red Hat
- Kubernetes Contributor
- CNCF Ambassador

# Agenda

Part 1: Containers & Docker

Part 2: Kubernetes

Part 3: Key Kubernetes Components

Part 4: Hands-On Demo

Part 5: Wrap-Up and Q&A

# Part 1: Containers & Docker

# Traditional Virtual Machines and Local Setups

Traditional virtual machines (VMs) virtualize an entire operating system, and running applications locally on a single OS can lead to conflicts.

Disadvantages:

- Heavy Resource Usage: VMs replicate the whole operating system, consuming substantial resources (RAM, CPU) on each VM instance.
- Lack of Consistency: Running apps locally can result in version conflicts (dependencies, environment variables)
- Inefficient Scaling: Spinning up a new VM is slower and requires more storage

# What are Containers?

Containers package an application with its dependencies, letting it run consistently across any environment.

Key Points:

- Containers occupy a smaller footprint compared traditional VMs.
- Containers allow consistency by packaging everything the application needs.
- Containers can be moved between environments without issues.

# Why Use Containers?

## Benefits:

- Portability: Containers can work across all computers and servers.
- Efficiency: Containers only carry what's needed for the app, making them smaller and faster than full virtual machines.
- Resource Isolation: Containers are independent, preventing conflicts.

# Introduction to Docker


Docker is the tool used to create and manage containers.

Key Commands:

`docker build`: Assembles the container image.

`docker run`: Runs the container.

`docker pull`: Downloads Docker images from the Docker Hub.





# Running a Container with Docker

Run an Nginx container:

```
docker run --name my-nginx -p 8080:80 nginx
```

Open a browser and visit localhost:8080 to see the Nginx welcome page.

# Part 2: Kubernetes

# Why Do We Need Kubernetes?

Complexity at Scale: Managing multiple containers across distributed environments is resource-intensive and complex.

- Scaling Needs: Adjusting container numbers manually to match traffic spikes or dips requires significant oversight.
- Configuration Drift: Without consistent orchestration, container configurations can drift across environments, causing reliability issues.

# Why Do We Need Kubernetes?

High Availability & Resilience:

- Fault Tolerance: Detecting and recovering from container or node failures promptly is crucial to avoid downtime.
- Service Discovery: Ensuring containers can locate each other (especially when they're dynamically scheduled across nodes) is challenging without orchestration.

# Why Do We Need Kubernetes?

Load Distribution:

- Traffic Management: Balancing incoming traffic across containers to avoid overload or resource wastage is difficult to manage manually.

# Kubernetes Overview

An orchestration platform that automates the deployment, scaling, and management of containerized applications.

Main Responsibilities:

- **Scaling:** Automatically adjusts the number of running containers based on real-time demand, adding or removing instances as needed.
- **Load Balancing:** Distributes incoming network traffic evenly across containers, ensuring efficient resource utilization and application performance.
- **Self-Healing:** Monitors container health, automatically replacing or restarting containers that fail, maintaining application availability and reliability.

# Key Kubernetes Concepts (Overview)

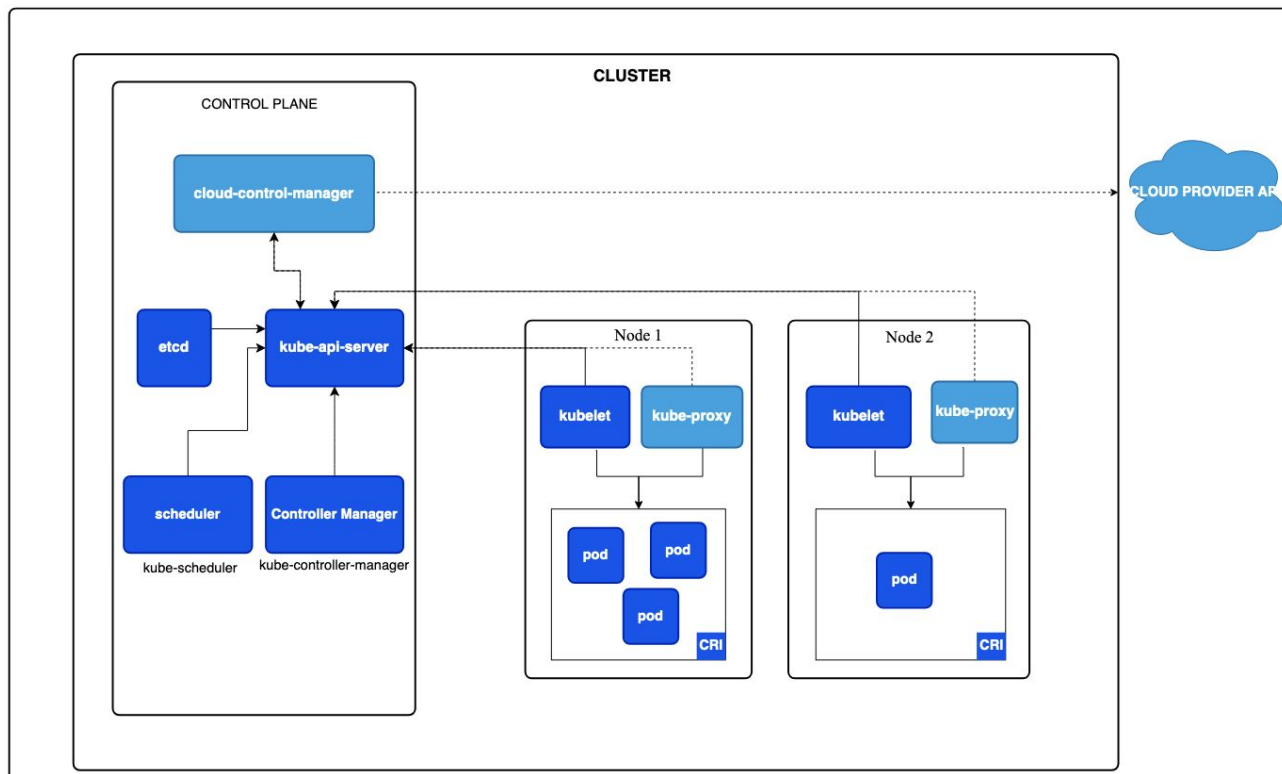
- Pod: The smallest deployable unit in Kubernetes, representing one or more containers that share storage, network, and a specification for how to run them.
- Deployment: A Kubernetes object that defines and manages the desired number of container replicas, providing scaling and update capabilities for applications.

# Key Kubernetes Concepts (Overview)

- Service: An abstraction that exposes a set of Pods as a network service, allowing stable access to the application regardless of dynamic Pod changes.
- Namespace: Kubernetes object that organizes and isolates resources, enabling multiple environments within the same Kubernetes cluster.



# Kubernetes Architecture



Pic credits: <https://kubernetes.io/docs/concepts/architecture/>

# Kubernetes Architecture

Control Plane: The core management layer of Kubernetes that maintains the desired state of the cluster.

- API Server: The central component that exposes the Kubernetes API and serves as the communication point for all Kubernetes operations.
- Scheduler: Assigns Pods to available nodes based on resource requirements and policies.
- Controller Manager: Ensures the actual state matches the desired state, handling tasks such as maintaining replicas and managing node lifecycles.
- etcd: A highly available key-value store that stores all cluster data, configurations, and states.

# Kubernetes Architecture

Worker Nodes: Nodes that run containerized applications and provide compute capacity to the cluster.

- Kubelet: An agent that manages Pod operations on each node, ensuring containers are running and healthy.
- Kube-proxy: A network proxy that manages network communication, routing traffic to the appropriate Pods within the cluster.

# Part 3: Key Kubernetes Concepts

# Pods - The Basic Unit

A Pod is the smallest deployable unit in Kubernetes, consisting of one or more tightly coupled containers that share storage, network, and a single IP address.

Key Features:

- **Shared Resources:** Containers within a Pod share network interfaces and storage volumes.
- **Use Case:** Often used to run a single container, but can contain multiple containers that need to work closely together.

*Example:* A web application Pod might contain both the main web server container and a sidecar container for logging or monitoring.

# Deployments

A Deployment is a Kubernetes object that defines and manages the desired state for a set of Pods, specifying the number of replicas, the container image version, and the update strategy.

Key Features:

- **Scaling:** Automatically increases or decreases the number of replicas based on demand.
- **Rolling Updates:** Ensures zero-downtime updates by gradually replacing old versions of Pods with new ones.
- **Self-Healing:** Automatically replaces failed Pods to maintain the specified replica count.

*Example:* A Deployment for a front-end application might specify three replicas, ensuring that three instances of the application are always available.

# Services

A Service is a Kubernetes resource that defines a stable IP and DNS name for a set of Pods, enabling communication between different parts of an application or external access to Pods.

Types of Services:

- ClusterIP: Exposes the Service within the cluster, allowing communication between internal Pods.
- NodePort: Exposes the Service on a static port on each node, allowing external access.
- LoadBalancer: Provisions an external load balancer, useful for cloud-based applications.

*Example:* A Service can expose a group of web server Pods to other Pods within the cluster or to external users, ensuring continuous access even if individual Pods are replaced or rescheduled.

# Part 4: Kubernetes in Action



# Demo

- Deploy a cluster using Kind
- Create an nginx webserver deployment
- Expose the deployment using a Service object
- Scale the Deployment
- Delete a pod and watch the deployment

# Part 5: Resources and Q/A

# Recap & Further Resources

- Containers: Provide consistent, portable application packaging by bundling code and dependencies. They enable reliable application deployment across diverse environments.
- Kubernetes: An orchestration platform that automates container management, ensuring scalability, fault tolerance, and efficient resource allocation for containerized applications.
  - Kind: A lightweight tool for running Kubernetes clusters locally, ideal for development, testing, and hands-on learning with Kubernetes.

# Recap & Further Resources

- Docker Tutorials: Resources for understanding Docker and containerization fundamentals – <https://www.docker.com/101-tutorial/>
- Kubernetes Documentation: In-depth guides on Kubernetes concepts, architecture, and components – [kubernetes.io/docs](https://kubernetes.io/docs)
- Kind Documentation: Learn more about setting up and managing local clusters – <https://kind.sigs.k8s.io/docs/user/quick-start/>



**THANK YOU :) :)**